

F7. Held-out predictions, pooled across all cutoffs. The model compresses the upper tail: no prediction exceeds 51.8% while observations reach 62.9%. The per-cutoff means reported in Table 5 are a different aggregation and do not coincide.

