

F15. The multi-layer perceptron used here: three hidden layers over 383 encoded inputs, 102,001 trainable parameters, ReLU activation, trained by stochastic gradient descent.

